

CMPE 200 Spring 2018

Quiz

Total Score: /20

Instructor: Sheng-Liang Song

Computer Engineering Department, San Jose State University

Quiz	Scores
1. <code>__sync_lock_test_and_set</code>	/2
2. Clock divide by 3	/2
3. Atomic APIs	/2
4. D/Q in Pipeline Register	/2
5. Pipeline Stall and Forwarding	/2
6. Extend the datapath to also support JALR	/2
7. 2-way set-associative cache	/2
8. Variable Size Page Table vs Hashed Page Table	/2
9. Virtual Address, TLB, Cache Tag (Physical Tag, Virtual Tag)	/2
10. Meltdown Paper	/2
Total	/20

Question 10: Read Meltdown Paper & Checkout and Run Examples.

<https://github.com/Shengliang/meltdown>

<https://drive.google.com/drive/u/1/folders/1lccGVN9Oj8fiy0wkUpSG3tzy4EI7RWBM>

Question 9: Virtual Address, TLB, Cache Tag (Physical Tag, Virtual Tag)

The following worksheet shows:

L1: is virtually tag cache

L2: is physical tag cache

TLB: direct-mapped, and fully associative

L1: direct-mapped and 2-way

L2: direct-mapped and 8-way (?)

<https://docs.google.com/spreadsheets/d/1hNXt-yOGN4I-nKOcEPXmLYcl9syXZ0XV7iO7XcJc2Uc/e/dit#gid=0>

<https://www.youtube.com/watch?v=YbE91A7dNAs>

Question 8: Variable Size Page Table vs Hashed Page Table

Given:

- 1) Page Size is 64KB
- 2) 32 Virtual Address Bus
- 3) Physical Memory 64GB

What is 1-level linear page table size per user space application?

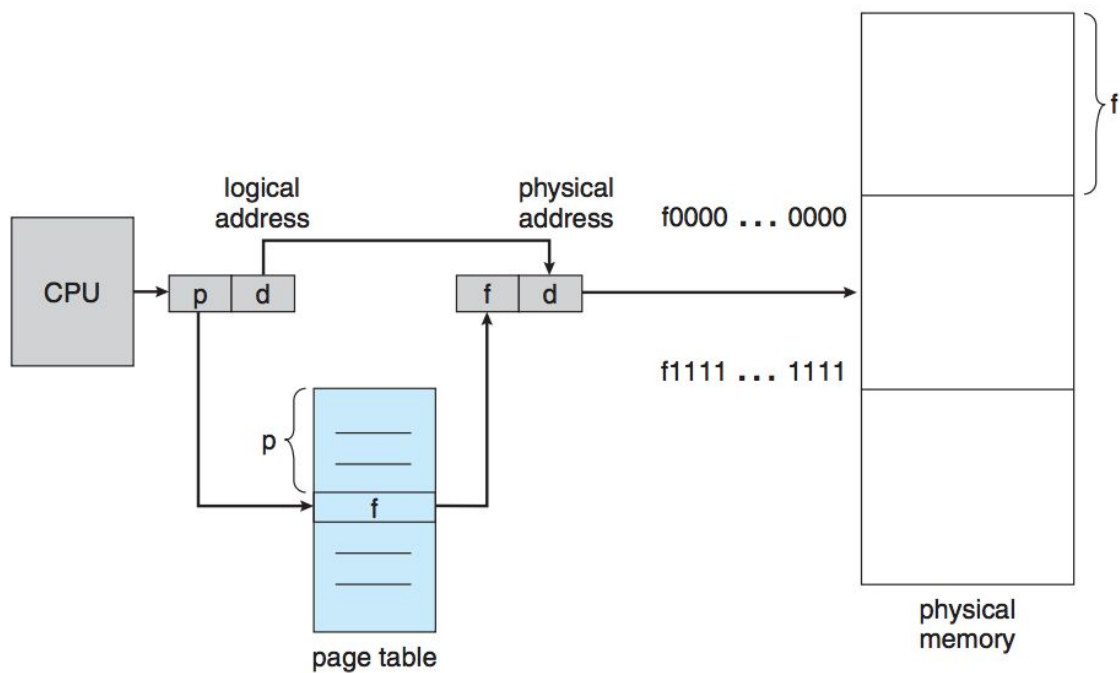


Figure 8.10 Paging hardware.

Page Size 64KB => d = 16 bit page offset

Virtual Page Number => p = 32 bit - 16 bit = 16 bits.

Page Table Entry (PTE): 4B

- A bit to indicate if a page exists (1 bit)
- PPN, physical page number for a memory-resident page (20 bits)
- Status bits for protection and usage (N bits)

64GB => 36 bit physical address

Physical Memory frame number (Physical Page Number) => f = 36 bit - 16 bit = **20 bits**

If a PTE is 32 bits (4 Byte), then

1-level linear page table size = $2^{16} * 4B = 64K * 4B = 256 KB$

// That can address $2^{16} * 64KB = 64K * 64KB = 2^{12}MB = 4GB$ memory

// or 32-bit virtual address can address 4GB memory.

What is 1-level linear page table size per user space application if the virtual address is 64-bit?

Virtual Page Number => p = 64 bit - 16 bit = 48 bits

1-level linear page table size = $2^{48} * \text{Page Table Entry Size} = 2^{48} * 4B = 2^{50} B = 1PB$. // That is not practical.

What is hashed page table size per user space application if the virtual address is 64-bit, and a hash function map 48-bits VPN to 16-bits hash-code?

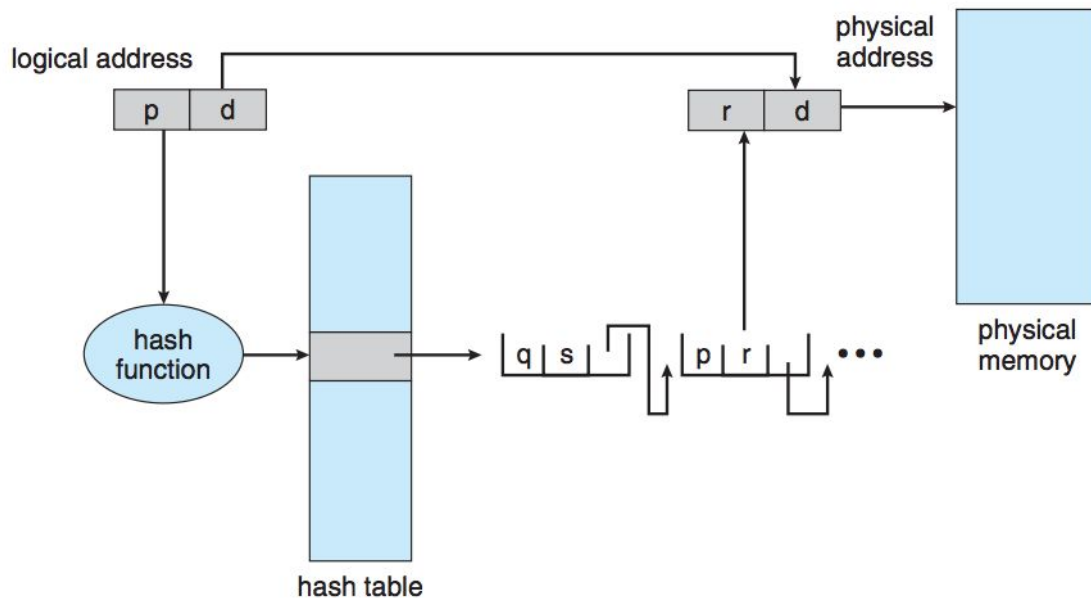


Figure 8.19 Hashed page table.

```
uint16_t get_hash_code(uint64_t vpn) {
    // Input: vpn = 0xFA44332211CE;
    // Output: hc = 0xFACE;
    return ((vpn >> 32) & 0xFF00) | (vpn & 0xFF);
}
```

```
uint64_t get_vpn_from_hash_code_and_tag(uint16_t hash_code, uint32_t tag) {
    // Input: hc = 0xFACE;
    // Input: tag = 0x44332211;
    // Output: vpn = 0xFA44332211CE;
    return ((hash_code & 0xFF00) << 32) | (tag << 8) | (hash_code & 0xFF);
}
```

Hashed Page Table Entry (PTE): = About 8 Bytes

- A bit to indicate if a page exists (1 bit)
- // VPN, virtual page number (48 bits)

- Virtual page number tag: **32 bits**
- Physical Memory frame number (Physical Page Number): $f = 36 \text{ bit} - 16 \text{ bit} = 20 \text{ bits}$
- Status bits for protection and usage (N bits)
- Next pointer (X bits) // Assume no next pointer.

Assume 1 hash entry saves 4 PTE. $8\text{B} * 4 = 32\text{B}$ (without next pointer)

[PTE] [PTE] [PTE] [PTE]

Hashed page table size = $2^{16} * 32\text{B} = 2^{21}\text{B} = 2 \text{ MB}$.

Figure 8.24 illustrates a x86 PAE system with 4-KB pages.




Figure 8.24 Page address extensions.

```
struct address_t {
    uint32_t page_directory_pointer_table_offset : 2;
    uint32_t page_directory_offset : 9;
    uint32_t page_table_offset : 9;
    uint32_t page_offset : 12;
};
```

Page Size: 4KB (12 bit page offset)

64GB => 36 bit physical address .

Physical Page Number size = 36 bit - 12 bit = 24 bits

- 1) Size of page directory pointer table: $4 * 4B = 16B$
- 2) Size of page directory: 512 entry ; $512 * 8B = 4KB$; Total Size: $4KB * 4 = 16KB$
- 3) Size of page table segment: 512 entry; $512 * 8B = 4KB$;
- 4) Size of page table: $4KB * 512 = 2MB$
- 5) Addressable memory from CR3: $4KB * 512 * 512 * 4 = 4GB$

Physical Address Extensions (PAE)

<http://www.superspeed.com/common/RdRp-Mem.pdf>

https://en.wikipedia.org/wiki/Physical_Address_Extension

<http://slideplayer.com/slide/7316543/>

<http://iips.icci.edu.iq/images/exam/Abraham-Silberschatz-Operating-System-Concepts---9th2012.12.pdf>

Question 7: Cache Questions

Question 5

- a) Let's run **your wrapper function** code on a five-stage pipeline architecture. We use a baseline register file where a newly written register value can be read in the subsequent cycle. Show the execution of above code by filling the timing diagram in the next page.

Assuming no forwarding logic.

- b) Repeat a. Now try to add hardware forwarding logic to eliminate stalls. (10 pts)
We still use a baseline register file where a newly written register value can be read in the subsequent cycle.

Please include the source and the destination of your forwarding logic.

With and without **branch delay slot** examples

- 1) Assume branch outcome is determined at EXE stage (in the book; assume **without branch delay slot**)

// if branch is taken; next PC = PC + 4 + offset = 104 + 4 + 200 = 308

100: lw \$7, 0(\$8)	F	D	E	M	W															
104: beqz \$7, +200		F	D	*	E	M	W													
108: addu \$6, \$7, \$0			F	*	D	E	M	W												
10C: addiu \$9, \$8, 4					F	D	E	M	W											
308: subi \$9, \$8, 4						F	D	E	M	W										

- 2) Assume branch outcome is determined at ID stage (if using gcc compiler; assume with **branch delay slot**)

// if branch is taken; next PC = PC + 4 + offset = 104 + 4 + 200 = 308 (assume 1 **branch delay slot**)

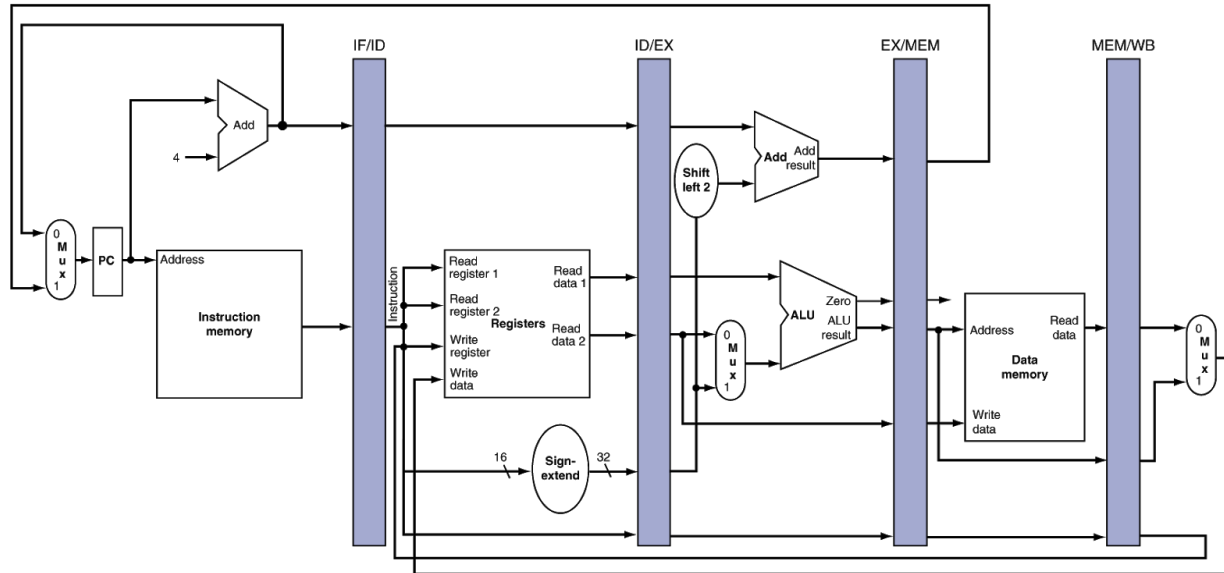
100:lw \$7, 0(\$8)		F	D	E	M	W													
104:beqz \$7, +200			F	*	*	D	E	M	W										
108:addu \$6,\$7,\$0						F	D	E	M	W		//branch delay slot: always execute							
308: subiu \$9,\$8,4							F	D	E	M	W								

// if branch is not taken, next pc = PC + 8 = 104 + 8 = 0x10C (assume with 1 **branch delay slot**)

100:lw \$7, 0(\$8)		F	D	E	M	W													
104:beqz \$7, +200			F	*	*	D	E	M	W										
108:addu \$6,\$7,\$0						F	D	E	M	W		//branch delay slot: always execute							
10C:addiu \$9,\$8,4							F	D	E	M	W								

Q4: The following diagram is a single-cycle pipeline diagram that we discussed.

add \$14, \$5, \$6	lw \$13, 24 (\$1)	add \$12, \$3, \$4	sub \$11, \$2, \$3	lw \$10, 20(\$1)
Instruction fetch	Instruction decode	Execution	Memory	Write-back



a.) Please please run your Q3 wrapper function code through the pipeline and decode all values (D and Q) in EXE/MEM pipeline register in the above diagram.

Pipeline Register	EXE/MEM.D	R	EXE/MEM.Q
Instruction Machine Code			
ISA Format			
opcode=?			
Next PC[31:0]			
ALU Result[31:0]			
RT Value[31:0]			
Dest Reg[4:0]			

CH2: Quiz

Please reference `list_tail_append()` in Question 1 and generate MIPS32 assembly and machine code for the following APIs. Your output should be similar to `list_tail_append()` output.

You may use gcc tools: `mips-mti-elf-gcc` and `mips-mti-elf-objdump`.

Reference: https://github.com/Shengliang/language/blob/master/c/atomic_api.c

Pick # from the following list:

N = Your ClassIndex % 14 + 1

Where N = {1,2,3,4,5,6,7,8,9,10,11,12,13,14}

<https://git.linaro.org/kernel/linux-linaro-stable.git>

1. type __sync_fetch_and_add (type *ptr, type value, ...)
2. type __sync_fetch_and_sub (type *ptr, type value, ...)
3. type __sync_fetch_and_or (type *ptr, type value, ...)
4. type __sync_fetch_and_and (type *ptr, type value, ...)
5. type __sync_fetch_and_xor (type *ptr, type value, ...)
6. type __sync_fetch_and_nand (type *ptr, type value, ...)
7. type __sync_add_and_fetch (type *ptr, type value, ...)
8. type __sync_sub_and_fetch (type *ptr, type value, ...)
9. type __sync_or_and_fetch (type *ptr, type value, ...)
10. type __sync_and_and_fetch (type *ptr, type value, ...)
11. type __sync_xor_and_fetch (type *ptr, type value, ...)
12. type __sync_nand_and_fetch (type *ptr, type value, ...)
13. bool __sync_bool_compare_and_swap (type *ptr, type oldval, type newval, ...)
14. type __sync_val_compare_and_swap (type *ptr, type oldval, type newval, ...)
15. type __sync_lock_test_and_set (type *ptr, type value, ...)
16. void __sync_lock_release (type *ptr, ...)

1) Create your github account

2) Add your google drive link and github link in [ClassIndex](#) Sheet


3) Go to <https://github.com/Shengliang/language>

4) Click fork


5) git clone https://github.com/<your_github_login>/language.git

Ex: <https://github.com/Shengliang/language.git>

6) https://github.com/Shengliang/language/blob/master/c/atomic_api.c




Appendix Review: Question 2 (1 points)



The above example assumes that t_{PLH} and t_{PHL} equal 20 ns for both AND and NOR gate. (50MHz : 20ns)

Reference:

<http://www.ece.ucsb.edu/Faculty/Johnson/ECE152A/handouts/L4%20-%20Propagation%20Delay,%20Circuit%20Timing%20&%20Adder%20Design.pdf>



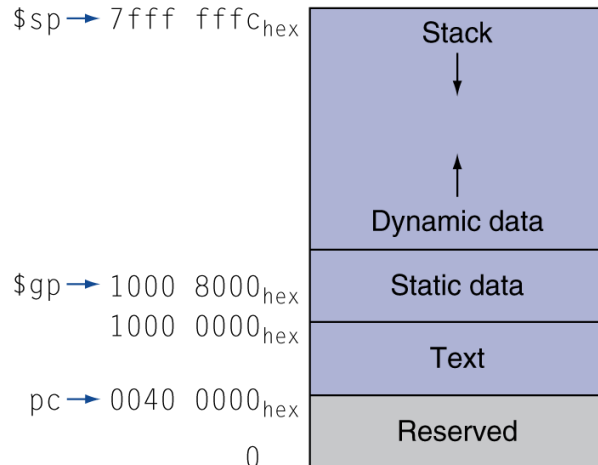
If clk_in is 1GHz, then what is the waveform for clk_out?

Assume “gate” delay: $t = (1/50) * \text{Period} = (1 \text{ ns})/50 = 1000\text{ps}/50 = 20\text{ps}$

Assume D-Flip-Flop delay: $T = 40\text{ps}$

CH2: Question 1 (1 points)

What are the printf output based on the following memory map? Please pick a address in the correct memory regions {Text, Static Data, Dynamic Data, Stack}.



```
#include<stdio.h>
#include<stdlib.h>
struct node_t {
    unsigned int value;
    struct node_t *next;
};

void print(struct node_t* n) {
    printf("print: %p\n", print);
    printf("&n: %p\n", &n);
    printf("&(n->value): %p\n", &(n->value));
    printf("&(n->next): %p\n", &(n->next));
    printf(" (n->next): %p\n", (n->next));
}

int main (void) {
    struct node_t* n = (struct node_t*) malloc(sizeof(struct node_t));
    n->value = 0xFACE;
    n->next = n;
    print(n);
    free(n);
    return 0;
}
```

I use the following command to generate objdump file.

```
mips-mti-elf-gcc -std=c99 -O3 list_append_v1_0.c
```

```
mips-mti-elf-objdump -D a.out > a.txt
```

```
20 void list_tail_append(struct list_t* list, struct node_t* n) {
21
22     // Check and acquire the lock, if available, else keep on checking
23     while (__sync_lock_test_and_set(&list->lock,1));
24
25     if (list->tail == NULL && list->head == NULL) {
26         list->head = n;
27         list->tail = n;
28     } else {
29         list->tail->next = n;
30         list->tail = n;
31     }
32
33     // Release the lock after execution of critical section
34     __sync_lock_release(&list->lock);
35 }
```

1) Which assembly line(s) are matching this GNU API ?

__sync_lock_release(&list->lock) (4 point)

2) Which assembly line(s) are matching this line C code? (6 point)

while (__sync_lock_test_and_set(&list->lock,1));

0040028c <list_tail_append>:

```
40028c: 24820008    addiu  v0,a0,8
400290: c0430000    ll     v1,0(v0)
400294: 24010001    li     at,1
400298: e0410000    sc     at,0(v0)
40029c: 1020fffc    beqz   at,400290 <list_tail_append+0x4>
4002a0: 00000000    nop
4002a4: 0000000f    sync
4002a8: 1460fff9    bnez   v1,400290 <list_tail_append+0x4>
4002ac: 00000000    nop
4002b0: 8c820004    lw     v0,4(a0)
4002b4: 10400007    beqz   v0,4002d4 <list_tail_append+0x48>
4002b8: 00000000    nop
4002bc: ac450004    sw     a1,4(v0)
4002c0: ac850004    sw     a1,4(a0)
4002c4: 0000000f    sync
4002c8: ac800008    sw     zero,8(a0)
4002cc: 03e00008    jr     ra
4002d0: 00000000    nop
4002d4: 8c830000    lw     v1,0(a0)
4002d8: 1460fff8    bnez   v1,4002bc <list_tail_append+0x30>
4002dc: 00000000    nop
4002e0: ac850000    sw     a1,0(a0)
4002e4: 1000fff7    b      4002c4 <list_tail_append+0x38>
4002e8: ac850004    sw     a1,4(a0)
```